

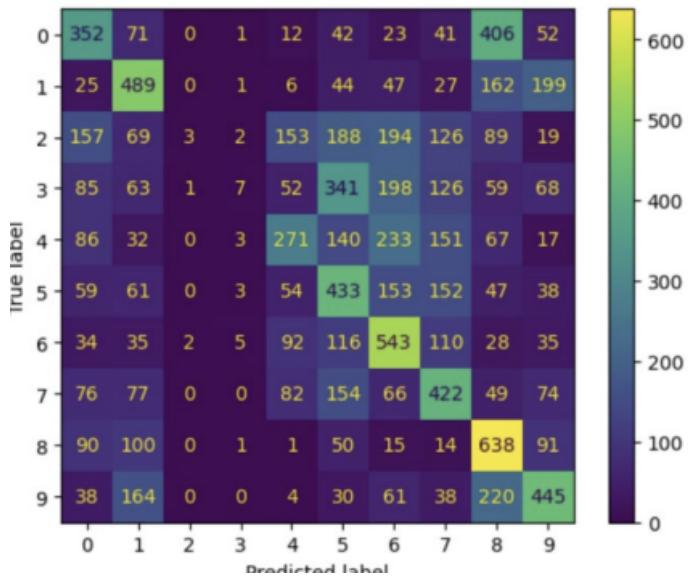
# 1. Image Classification with CNN (PyTorch Lightning)

## Model Overview

- **Input:** A dataset consisting of images from **10 categories**—plane, car, bird, cat, deer, dog, frog, horse, ship, and truck.
- **Output:** The predicted class label for each image and its **confidence score** (probability).

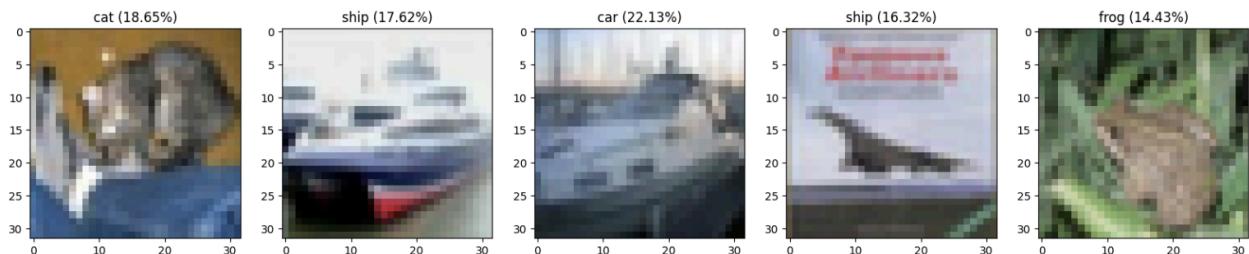
Layer (type:depth-idx)	Output Shape	Param #
CNN	[32, 10]	--
-Conv2d: 1-1	[32, 6, 28, 28]	456
-MaxPool2d: 1-2	[32, 6, 14, 14]	--
-Conv2d: 1-3	[32, 16, 10, 10]	2,416
-MaxPool2d: 1-4	[32, 16, 5, 5]	--
-Linear: 1-5	[32, 120]	48,120
-Linear: 1-6	[32, 84]	10,164
-Linear: 1-7	[32, 10]	850
-Softmax: 1-8	[32, 10]	--
<hr/>		
Total params:	62,006	
Trainable params:	62,006	
Non-trainable params:	0	
Total mult-adds (Units.MEGABYTES):	21.06	
<hr/>		
Input size (MB):	0.39	
Forward/backward pass size (MB):	1.67	
Params size (MB):	0.25	
Estimated Total Size (MB):	2.31	
<hr/>		

- **Hardware Requirements:** The model primarily runs on **GPU (CUDA:0)** for faster computation. If a GPU is unavailable, the code automatically switches to the **CPU**.
- **Data Statistics:** The model evaluation includes **precision, recall, and F1-score** as key performance metrics.
- **Learning Curve:** The model tracks accuracy and loss throughout training and validation to visualize learning progression.
- **Metrics:** **Softmax activation** is applied in the final layer to compute class probabilities and determine the most likely label.
- **Demo Result:** Includes a statistical evaluation of the model's performance and sample predictions to illustrate classification accuracy.



## Key Functions & Features (Cheat Sheet)

- **Dataset:**
  - The dataset consists of **60,000 images** divided into:
    - **40,000** for training
    - **10,000** for validation
    - **10,000** for testing
- **Training Process:**
  - Utilizes **PyTorch Lightning Trainer**
  - Runs for a **maximum of 1 epoch**
  - Records **training and validation loss & Accuracy**
- **Evaluation & Testing:**
  - The trained model is tested on the **test dataset**
  - Results are assessed using:
    - **Classification Precision**
    - **Confusion Matrix** (to visualize class-wise accuracy)
- **Final Output:**
  - The model's predictions are demonstrated using the **first 5 images** from the test set to visually represent its classification capabilities.



## 2. Image Classification with EfficientNetV2s (PyTorch Lightning)

### Model Overview

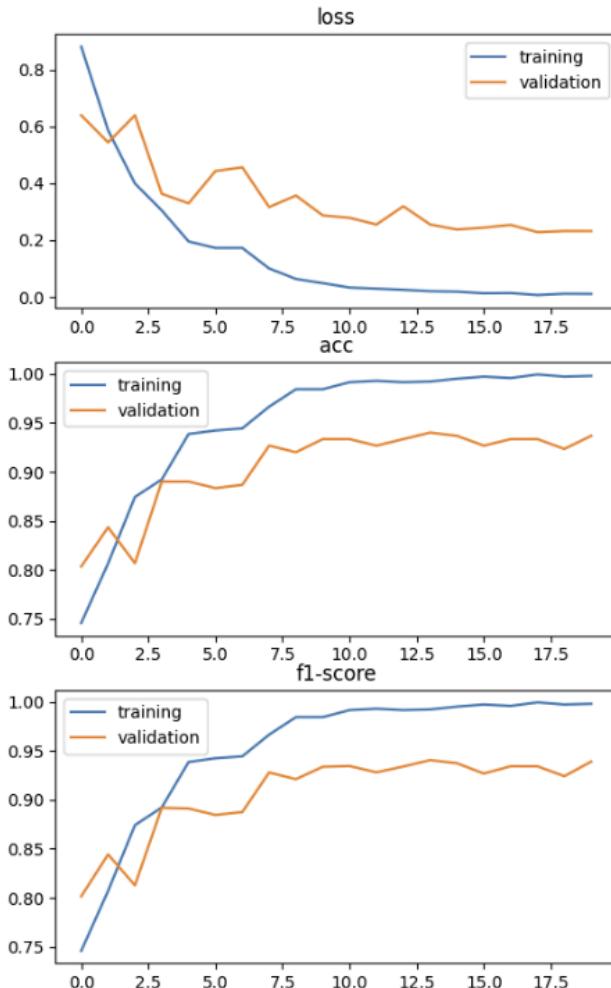
- **Input:** A dataset containing images of **10 animal classes**—butterfly, cat, chicken, cow, dog, elephant, horse, sheep, spider, and squirrel.
- **Output:** The predicted class for each image and a **confidence score** (probability).

Layer (type)	Output Shape	Param #
Conv2d-1	[64, 24, 112, 112]	648
BatchNorm2d-2	[64, 24, 112, 112]	48
SiLU-3	[64, 24, 112, 112]	0
Conv2d-4	[64, 24, 112, 112]	5,184
BatchNorm2d-5	[64, 24, 112, 112]	48
SiLU-6	[64, 24, 112, 112]	0
StochasticDepth-7	[64, 24, 112, 112]	0
FusedMBConv-8	[64, 24, 112, 112]	0
Conv2d-9	[64, 24, 112, 112]	5,184
BatchNorm2d-10	[64, 24, 112, 112]	48
SiLU-11	[64, 24, 112, 112]	0
StochasticDepth-12	[64, 24, 112, 112]	0
FusedMBConv-13	[64, 24, 112, 112]	0
Conv2d-14	[64, 96, 56, 56]	20,736
BatchNorm2d-15	[64, 96, 56, 56]	192
SiLU-16	[64, 96, 56, 56]	0
Conv2d-17	[64, 48, 56, 56]	4,608
BatchNorm2d-18	[64, 48, 56, 56]	96
FusedMBConv-19	[64, 48, 56, 56]	0
Conv2d-20	[64, 192, 56, 56]	82,944
BatchNorm2d-21	[64, 192, 56, 56]	384
SiLU-22	[64, 192, 56, 56]	0
...		
Forward/backward pass size (MB):	20629.03	
Params size (MB):	77.02	
Estimated Total Size (MB):	20742.80	

- **Hardware Requirements:** The model primarily runs on **CUDA:0 (GPU)** for efficient training. If a GPU is unavailable, it falls back to the **CPU**. The model used for classification is **LitEfficientNetV2**.
- **Data Statistics:**
  - Total **2000 images**, split into:
    - **1400** for training

- 300 for validation
- 300 for testing

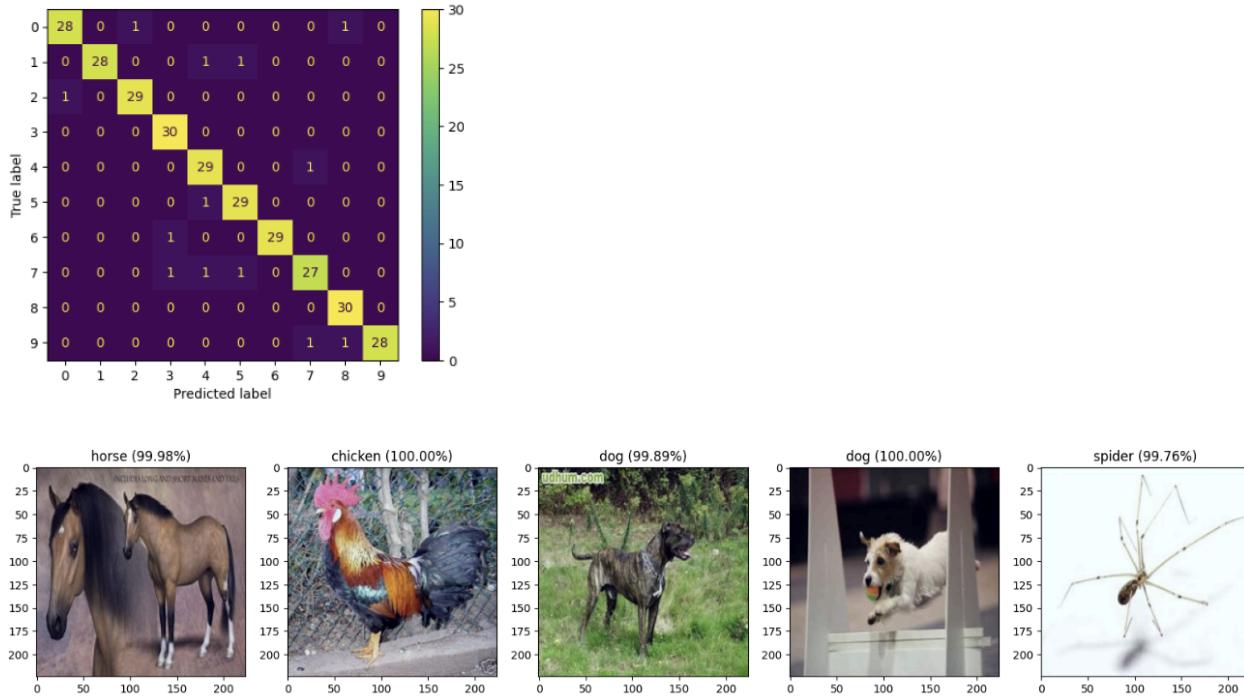
- **Learning Curve:** Tracks accuracy and loss progression across training and validation phases.



- **Metrics:**
  - **Softmax activation** is applied in the final layer to compute class probabilities.
  - Model performance is assessed using **precision, recall, and F1-score** to measure classification accuracy.
- **Demo Result:**
  - The model's evaluation statistics, including the **confusion matrix**, are displayed.
  - A usage demo illustrates the model's predictions, highlighting cases where it correctly classifies images with high confidence.
  - Incorrect classifications are observed, particularly in cases where confidence scores are low.

## Key Functions & Features (Cheat Sheet)

- **Data Processing:**
  - The input images are **downloaded and randomized** to include different orientations.
  - The dataset consists of **2000 images**, divided into:
    - **1400** for training
    - **300** for validation
    - **300** for testing
- **Training & Evaluation:**
  - The training dataset is used to **train and save the model** for future use.
  - Performance is visualized using a **confusion matrix** to analyze class-wise accuracy.
  - Model predictions are displayed, showing the **confidence percentage** for each classification.



## 3-2. Object Detection with YOLOv8 (Basic Script)

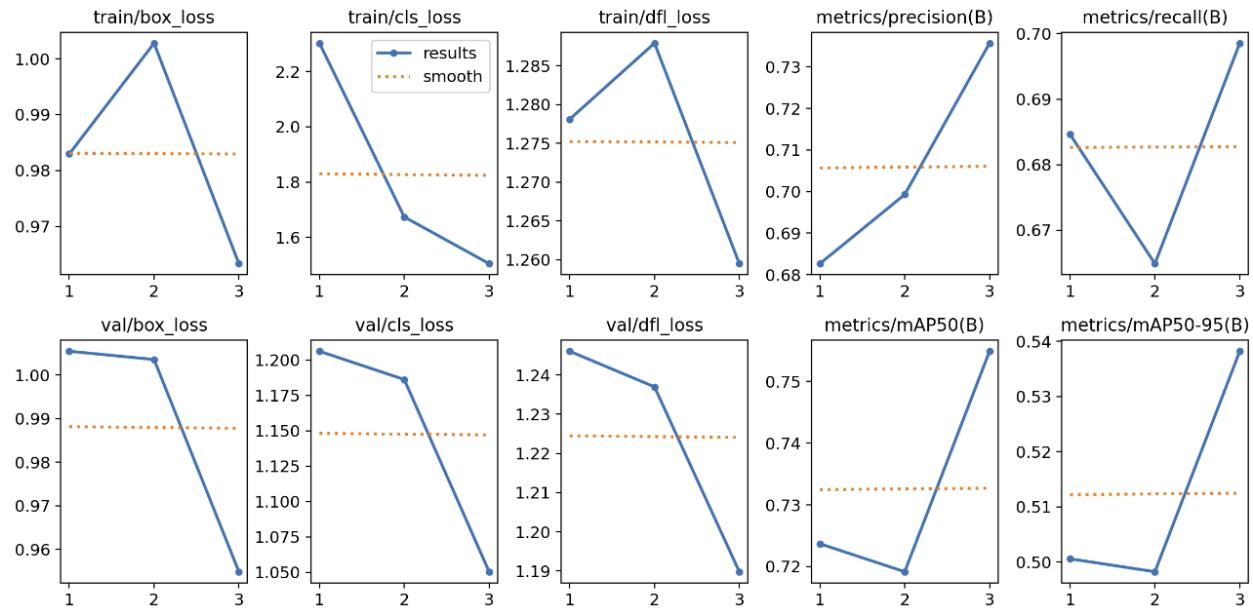
### Model Overview

- **Input:** A dataset containing images with **20 object categories**, including **person, bird, cat, dog, horse, sheep, airplane, bicycle, boat, bus, car, motorbike, train, bottle, chair, dining table, potted plant, sofa, and TV monitor**.

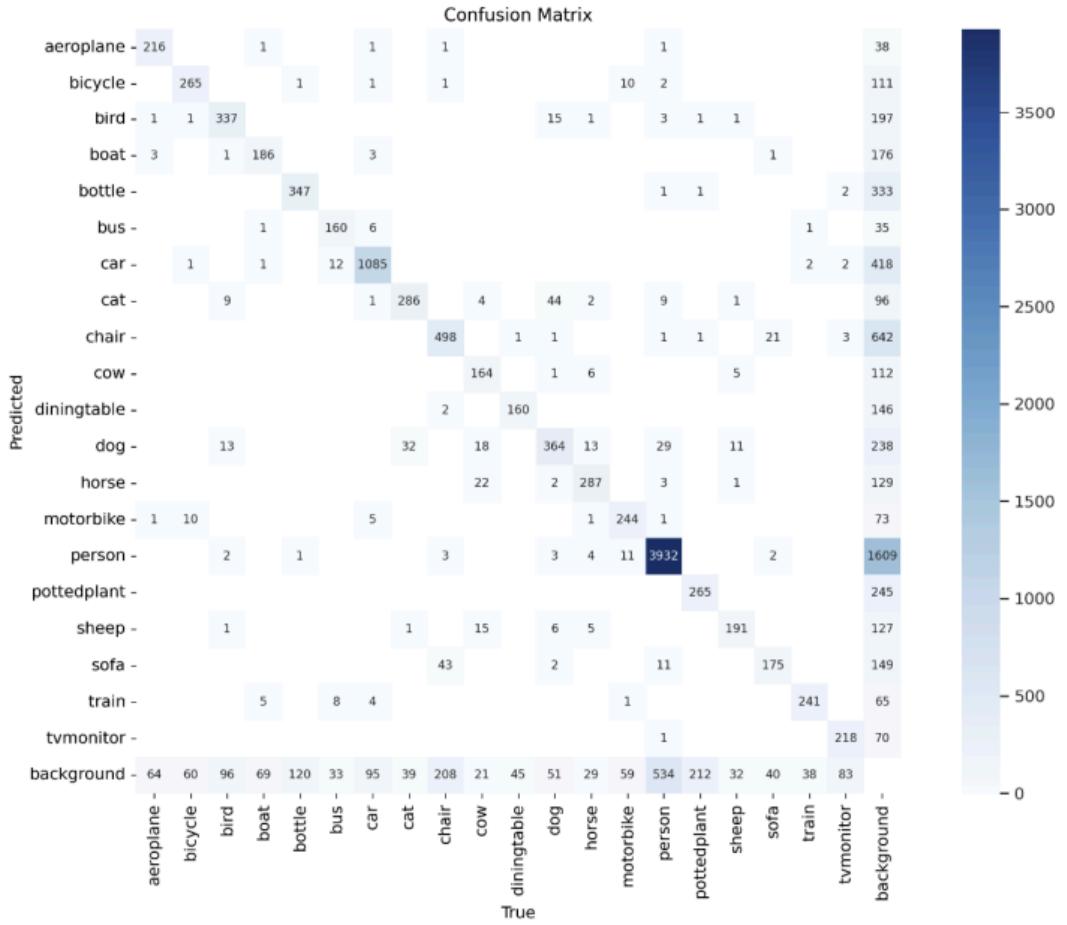
- **Output:** The model detects and classifies **multiple objects** in an image, assigning each **object class** along with a **confidence score**.

	from	n	params	module	arguments
0	-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1	-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2	-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4	-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8	-1	1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9	-1	1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	37248	ultralytics.nn.modules.block.C2f	[192, 64, 1]
16	-1	1	36992	ultralytics.nn.modules.conv.Conv	[64, 64, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	123648	ultralytics.nn.modules.block.C2f	[192, 128, 1]
19	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]
...					

- **Hardware Requirements:**
  - The model is designed to run on **CUDA (Nvidia GPU)** for efficient computation.
  - **YOLOv8** is used as the object detection framework.
  - The training process consists of **3 epochs**.
- **Data Statistics:**
  - Each training image is **resized to 640x640 pixels** to standardize the input.
  - The dataset contains **16,551 images**, providing a large sample size for improved model performance.
- **Learning Curve:**
  - The model learns to detect objects by iterating over the dataset during training.



- **Metrics:**
  - The model's performance is validated using the **VOC validation dataset (VOC val)**.
- **Demo Result:**
  - A confusion matrix is generated to visualize the model's performance in classifying detected objects.
  - The model correctly detects and classifies most objects, though some misclassifications occur.



## Key Functions & Features (Cheat Sheet)

- **Object Detection vs. Image Classification:**
  - Unlike the previous two models, which only **classify whole images**, this model can **detect and classify multiple objects within an image**, identifying their locations.
- **Training Process:**
  - **16,551 images** are used for training, significantly improving the model's accuracy.
  - Training images are resized to **640x640 pixels** before being fed into the model.
- **Evaluation:**
  - The confusion matrix illustrates how accurately the model assigns object classes.
  - The VOC validation dataset is used to measure performance.
  -

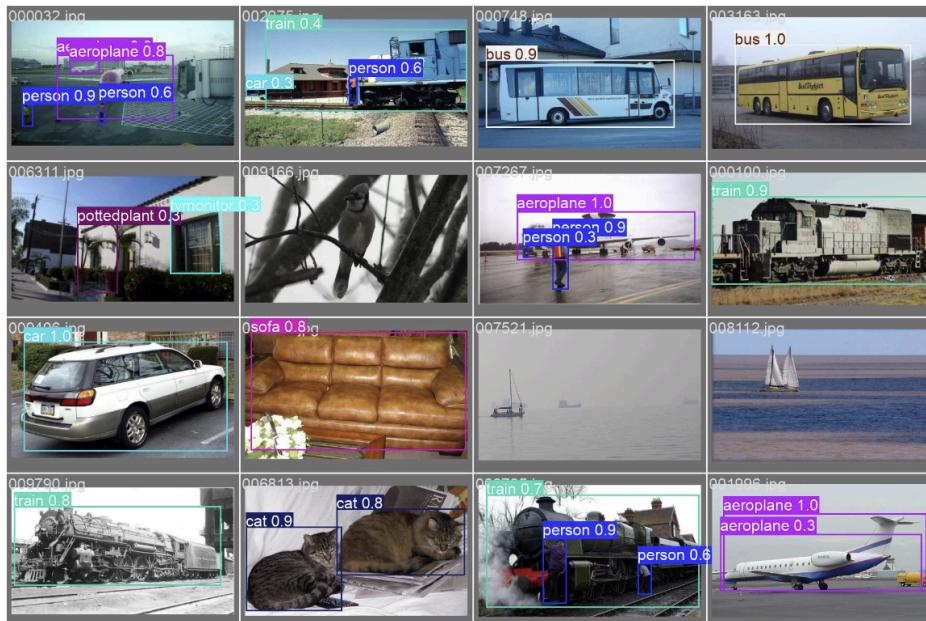
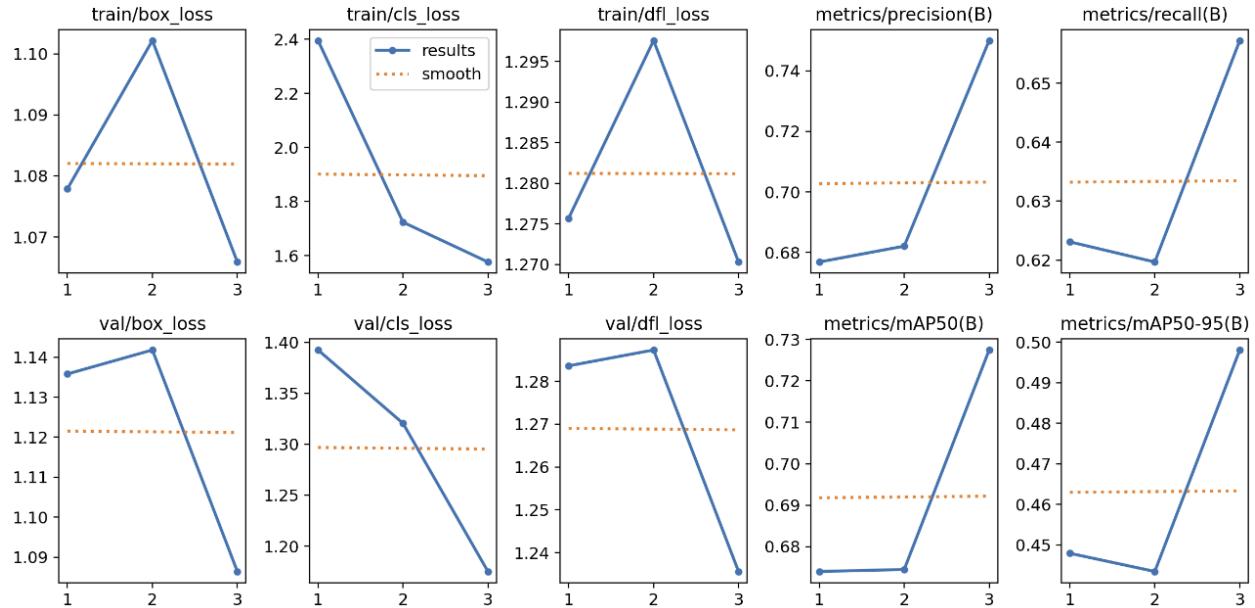
### 3-3. Object Detection Using YOLOv8 (Custom Dataset)

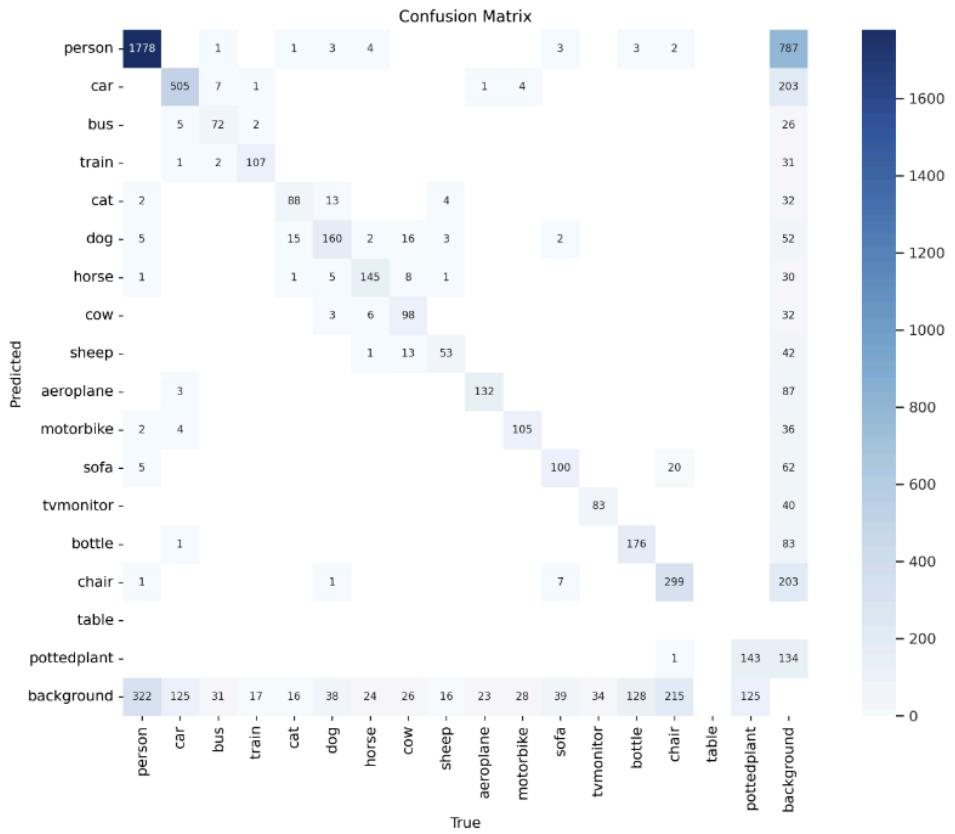
#### Model Overview

- **Input:** A dataset consisting of images containing **20 object categories**, including **person, bird, cat, dog, horse, sheep, airplane, bicycle, boat, bus, car, motorbike, train, bottle, chair, dining table, potted plant, sofa, and TV monitor**.
- **Output:** The model identifies **objects within an image**, assigning them to specific classes along with a **confidence score**.

	from	n	params	module	arguments
0	-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1	-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2	-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4	-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8	-1	1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9	-1	1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	37248	ultralytics.nn.modules.block.C2f	[192, 64, 1]
16	-1	1	36992	ultralytics.nn.modules.conv.Conv	[64, 64, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	123648	ultralytics.nn.modules.block.C2f	[192, 128, 1]
19	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]

- **Hardware Requirements:**
  - The model runs on **CUDA:0**, utilizing an **Nvidia GPU** for efficient processing.
- **Data Statistics:**
  - The dataset is sourced from a **publicly available dataset**, as referenced in the code.
  - **20% of the dataset** is allocated for testing.
- **Learning Curve & Demo Results:**
  - The model's performance is evaluated through detection accuracy and object classification confidence.





## Key Features & Function Overview (Cheat Sheet)

- **Object Detection vs. Image Classification:**
  - Unlike previous models, which solely **classify entire images**, this model can **detect multiple objects within an image** while also classifying them.
- **Public Dataset Utilization:**
  - The model is trained using a **public dataset**, making it adaptable to specific object detection tasks.
  - **20% of the dataset is reserved for testing**, ensuring reliable performance evaluation.
  -

## 4. Semantic Segmentation Using DeepLabV3 (PyTorch Lightning)

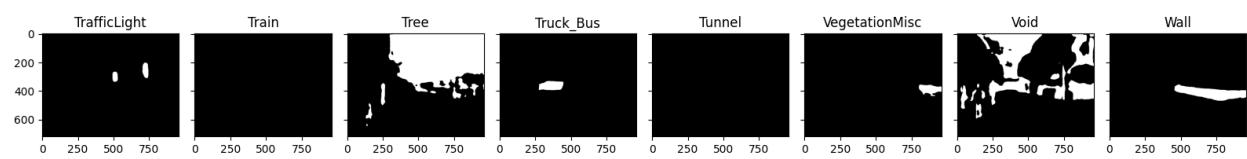
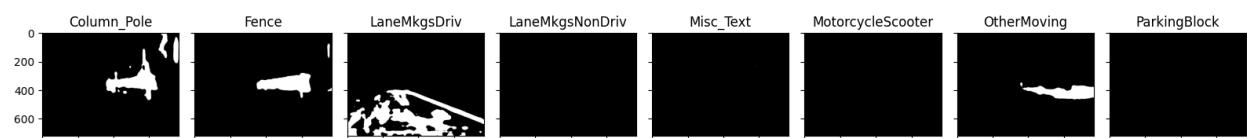
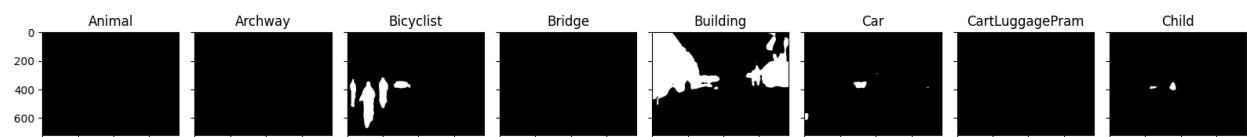
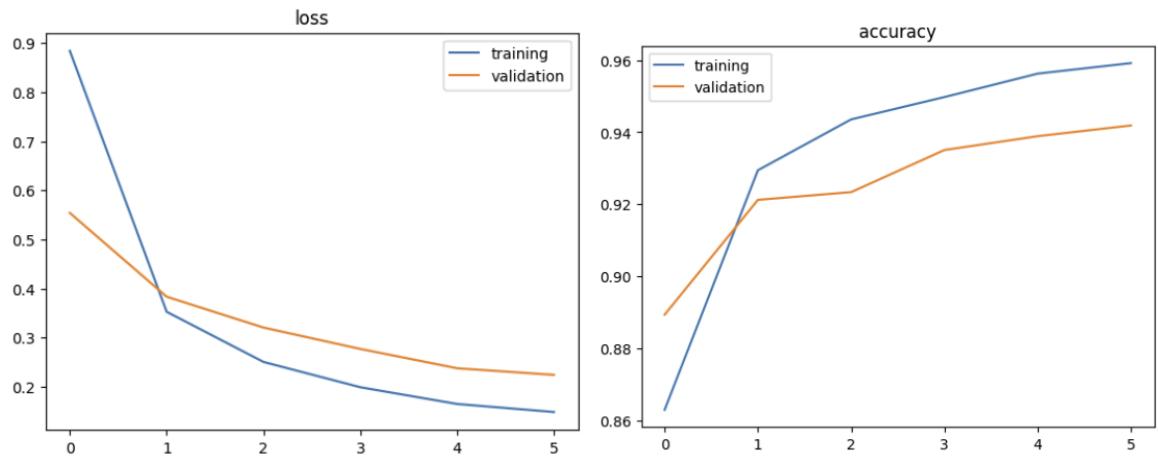
### Model Overview

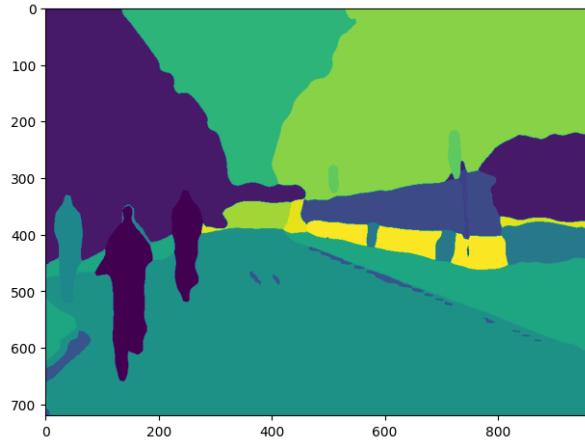
- **Input:**
  - A dataset containing **101 image pairs** with dimensions **960x720 pixels**.
  - Each mask file is labeled with an "**\_L**" **suffix** in the filename.
  - All images and ground truth masks are in **uncompressed 24-bit color PNG format**.

- **Output:**
  - **Segmented images**, where each object in an image is assigned a distinct region.

```
DeepLabV3(
    backbone): IntermediateLayerGetter(
        (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
        (layer1): Sequential(
            (0): Bottleneck(
                (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
                (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
                (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
                (relu): ReLU(inplace=True)
                (downsample): Sequential(
                    (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
                )
            )
            (1): Bottleneck(
                (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
                (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            ...
            (2): ReLU()
            (3): Dropout(p=0.1, inplace=False)
            (4): Conv2d(256, 21, kernel_size=(1, 1), stride=(1, 1))
        )
    )
)
```

- **Hardware Requirements:**
  - The model runs on **CUDA:0** by default but can fall back to the **CPU** if necessary.
- **Data Statistics:**
  - Dataset split: **70 images for training, 16 for validation, and 15 for testing.**
- **Learning Curve & Demo Results:**
  - The model is assessed by observing its segmentation accuracy and how well it distinguishes different objects within images.





## Key Features & Function Overview (Cheat Sheet)

- **DeepLabV3 Model Usage:**
  - The DeepLabV3 model is trained to **segment images into multiple regions** based on the objects present.
- **Segmentation Output:**
  - The model's output consists of **images segmented into distinct parts**, highlighting the locations of different objects.
- **Performance Evaluation:**
  - The quality of segmentation is assessed visually by comparing the **predicted segmentations** with **ground truth masks**.