

Messaging with RabbitMQ

<https://www.rabbitmq.com/#getstarted>

About RabbitMQ

RabbitMQ is an open-source distributed message/streaming broker.

RabbitMQ supports AMQP (Advanced Messaging Queuing Protocol), MQTT (Message Queue Telemetry Transport protocol)

AMQP, MQTT are an open standard application layer protocols.

RabbitMQ accepts, stores, and forwards messages: blobs of data

Jargons

<p>Producing</p> <p>A program which sends message is a <i>producer</i>.</p>	
<p>Queue : a post box living inside RabbitMQ.</p>	<p>queue_name</p> 
<p>Consuming</p> <p>A <i>consumer</i> is a program that mostly waits to receive messages.</p>	

RabbitMQ is a push model.

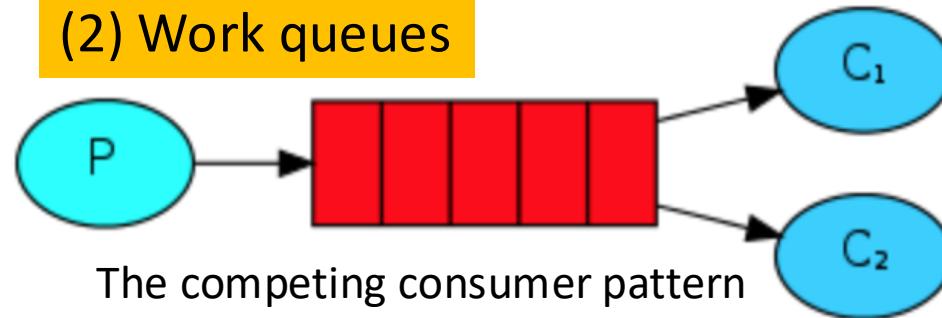
Queue sends message to consumers



(1) Hello world

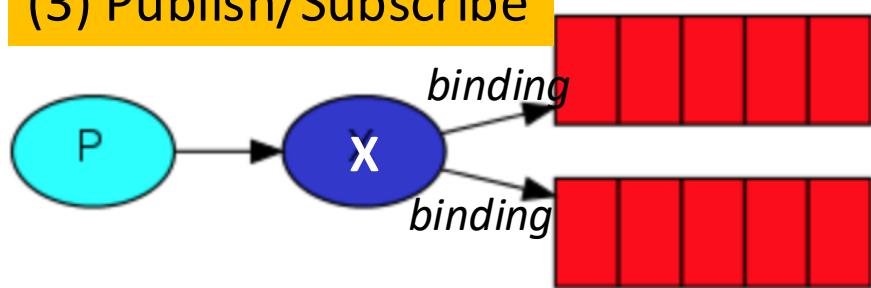


(2) Work queues



(7) Publisher confirms

(3) Publish/Subscribe



Exchange type: direct, topic, headers, fanout

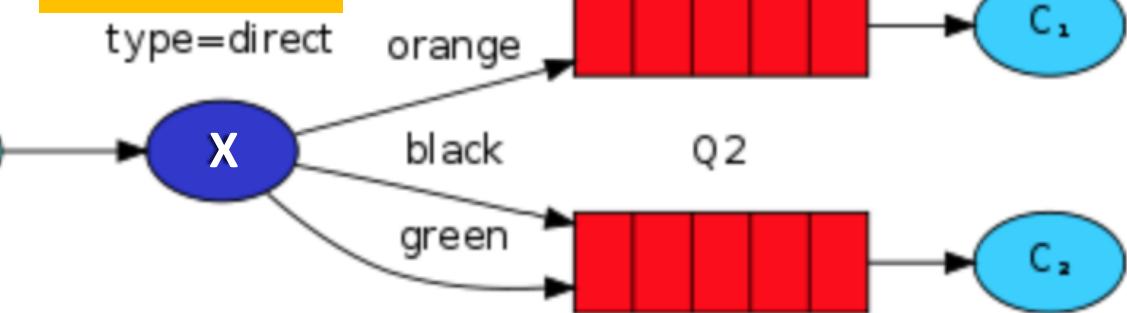
(5) Topics

* Can substitute for exactly one word

Can substitute for zero or more words

(4) Routing

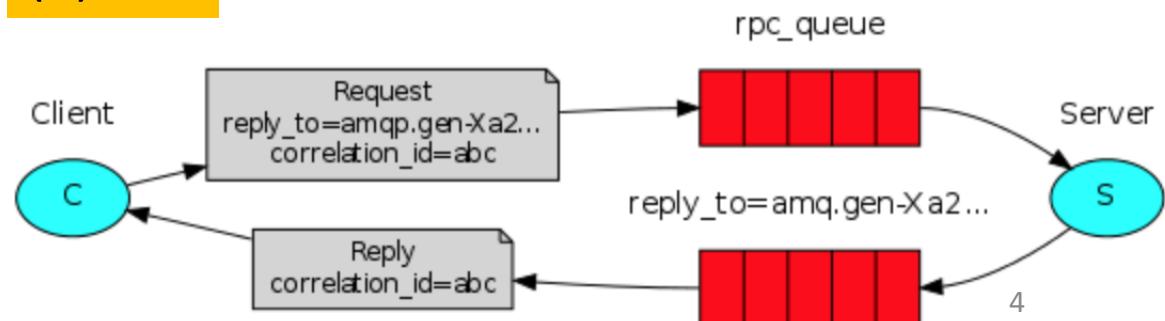
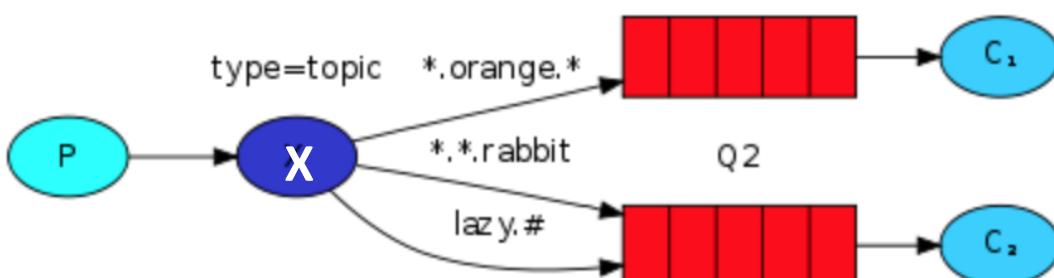
type=direct



Exchange type: direct, topic, headers, fanout

binding key == routing key

(6) RPC



Pre-requisite

Install RabbitMQ on local machine.

RabbitMQ server

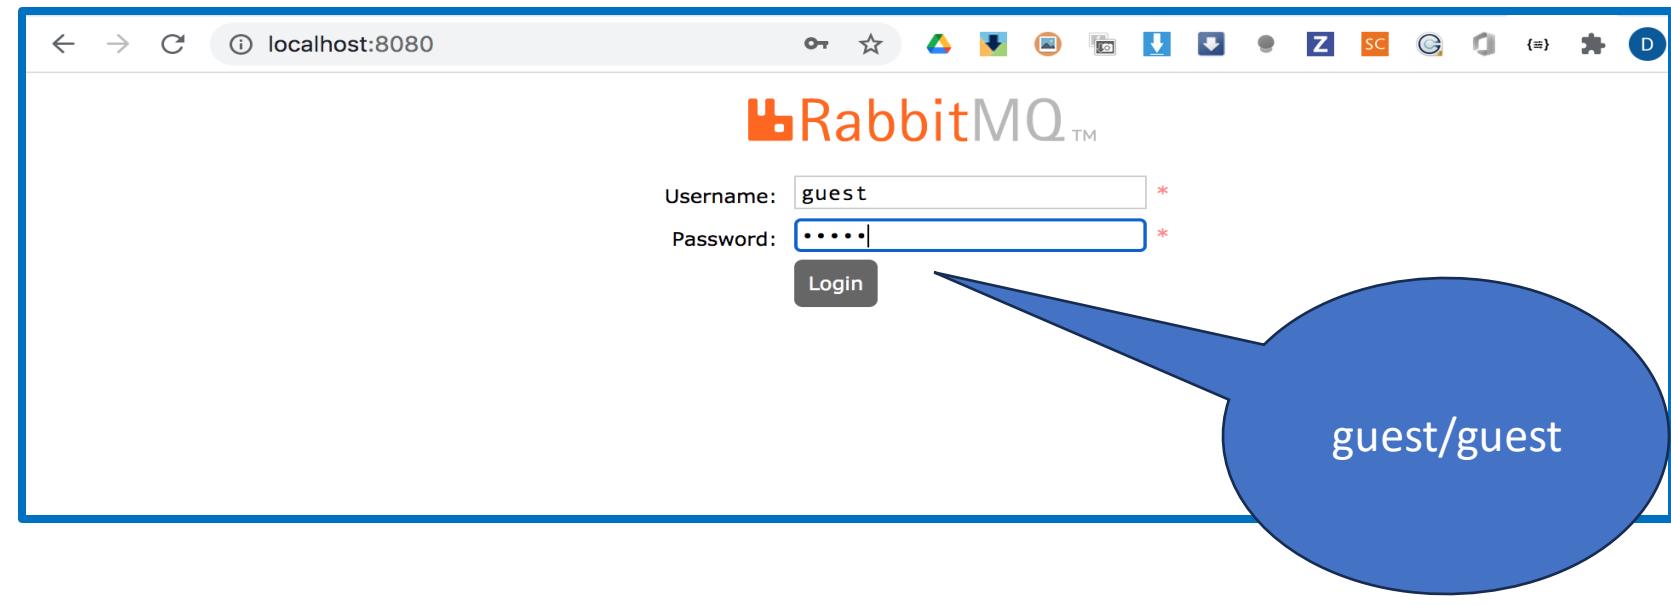
```
docker run --name my-rabbitmq -p 5672:5672 -d rabbitmq:4.1.4
```

RabbitMQ management

```
docker run --name my-rabbitmq-management -p 8080:15672 -d rabbitmq:4.1.4-management
```

<https://hub.docker.com/rabbitmq>

Pre-requisite



<http://localhost:8080>

```
docker exec my-rabbitmq-management rabbitmqctl status  
docker exec my-rabbitmq-management rabbitmqctl
```

<https://hub.docker.com/rabbitmq>

RabbitMQ Management

RabbitMQ™ RabbitMQ 3.9.5 Erlang 24.0.5

Refreshed 2021-09-19 00:38:04 Refresh every 5 seconds ▾

Virtual host All ▾

Cluster **rabbit@my-rabbit**

User **guest** Log out

Overview Connections Channels Exchanges Queues Admin

Overview

▼ Totals

Queued messages last minute ?

Currently idle

Message rates last minute ?

Currently idle

Global counts ?

Connections: 0 Channels: 0 Exchanges: 7 Queues: 0 Consumers: 0

▼ Nodes

Name	File descriptors ?	Socket descriptors ?	Erlang processes	Memory ?	Disk space	Uptime	Info	Reset stats	+
rabbit@my-rabbit	37 1048576 available	0 943629 available	393 1048576 available	140 MiB 796 MiB high watermark	32 GiB 48 MiB low watermark	4m 9s	basic disc 2 rss	This node All nodes	

► Churn statistics

► Ports and contexts

► Export definitions

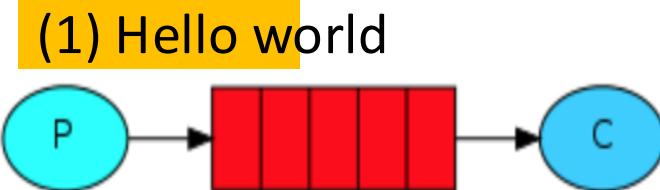
► Import definitions

[HTTP API](#) [Server Docs](#) [Tutorials](#) [Community Support](#) [Community Slack](#) [Commercial Support](#) [Plugins](#) [GitHub](#) [Changelog](#)

Python library as RabbitMQ client

- `python -m pip install pika --upgrade`

(1) Hello world

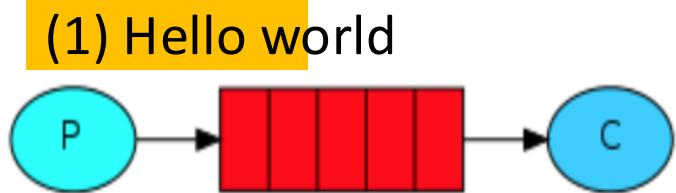


python 1_HelloReceive.py
python 1_HelloSent.py

Producer (P)

```
1 import sys
2
3 #!/usr/bin/env python
4 import pika
5
6 connection = pika.BlockingConnection(pika.ConnectionParameters('localhost'))
7 channel = connection.channel()
8 channel.queue_declare(queue='hello')
9
10 channel.basic_publish(exchange='',
11                       routing_key='hello',
12                       body='Hello World!')
13 print("[x] Sent 'Hello World!'")
14 connection.close()
```

(1) Hello world

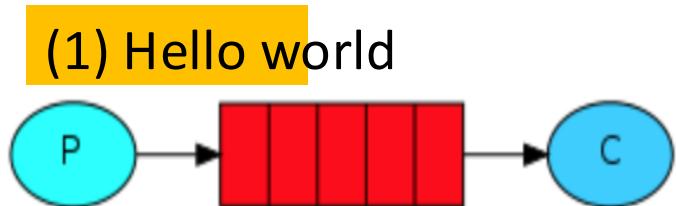


python 1_HelloReceive.py
python 1_HelloSent.py

Consumer (C)

```
1  #!/usr/bin/env python
2  import pika, sys, os
3
4  def main():
5      connection = pika.BlockingConnection(pika.ConnectionParameters(host='localhost'))
6      channel = connection.channel()
7
8      channel.queue_declare(queue='hello')
9
10     def callback(ch, method, properties, body):
11         print(" [x] Received %r" % body)
12
13     channel.basic_consume(queue='hello', on_message_callback=callback, auto_ack=True)
14
15     print(' [*] Waiting for messages. To exit press CTRL+C')
16     channel.start_consuming()
17
18 if __name__ == '__main__':
19     try:
20         main()
21     except KeyboardInterrupt:
22         print('Interrupted')
23         try:
24             sys.exit(0)
25         except SystemExit:
26             os._exit(0)
```

(1) Hello world



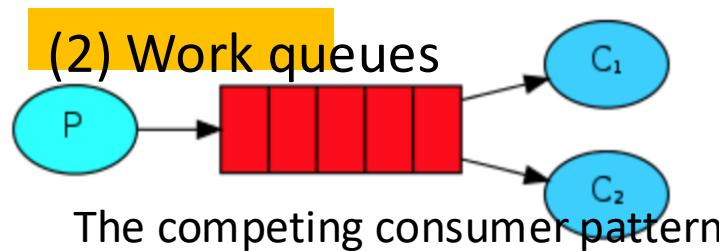
```
python 1_HelloReceive.py  
python 1_HelloSend.py
```

DEMO

(2) Work queue/ Task queue

The assumption behind a work queue is that each task is delivered to exactly one worker.

Round robin dispatching



The competing consumer pattern

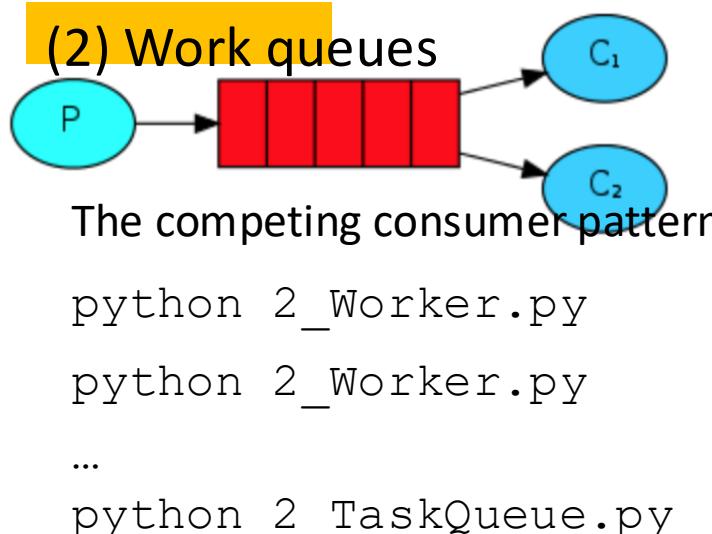
`python 2_Worker.py`

`python 2_Worker.py`

...

`python 2_TaskQueue.py`

(2) Work queue/ Task queue (cont.)



Producer (P)

```
1 #!/usr/bin/env python
2 import pika
3 import sys
4
5 connection = pika.BlockingConnection(
6     pika.ConnectionParameters(host='localhost'))
7 channel = connection.channel()
8
9 channel.queue_declare(queue='task_queue', durable=True)
10
11 message = ' '.join(sys.argv[1:]) or "Hello World!"
12 channel.basic_publish(
13     exchange='',
14     routing_key='task_queue',
15     body=message,
16     properties=pika.BasicProperties(
17         delivery_mode=2, # make message persistent
18     ))
19 print("[x] Sent %r" % message)
20 connection.close()
```

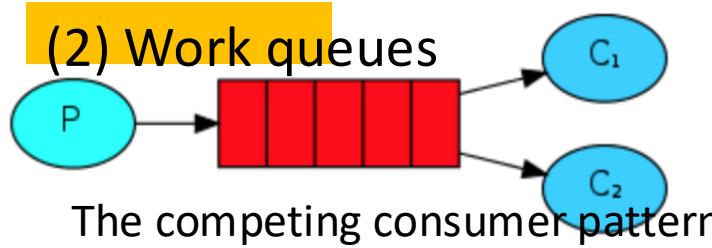
Message durability

To ensure if RabbitMQ quits or crashes, it won't forget queues

#make queue persistent

Message will be stored to disk.

(2) Work queue/ Task queue (cont.)



(2) Work queues
The competing consumer pattern

python 2_Worker.py

python 2_Worker.py

...

python 2_TaskQueue.py

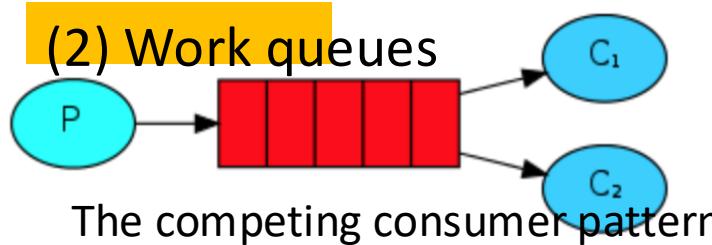
Consumer (C)

Message durability

```
1  #!/usr/bin/env python
2  import pika
3  import time
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.queue_declare(queue='task_queue', durable=True)
10 print(' [*] Waiting for messages. To exit press CTRL+C')
11
12
13 def callback(ch, method, properties, body):
14     print(" [x] Received %r" % body.decode())
15     time.sleep(body.count(b'.'))
16     print(" [x] Done")
17     ch.basic_ack(delivery_tag=method.delivery_tag)
18
19
20 channel.basic_qos(prefetch_count=1)
21 channel.basic_consume(queue='task_queue', on_message_callback=callback)
22
23 channel.start_consuming()
```

To ensure if RabbitMQ quits or crashes, it won't forget queues

(2) Work queue/ Task queue (cont.)



python 2_Worker.py

python 2_Worker.py

...

python 2_TaskQueue.py

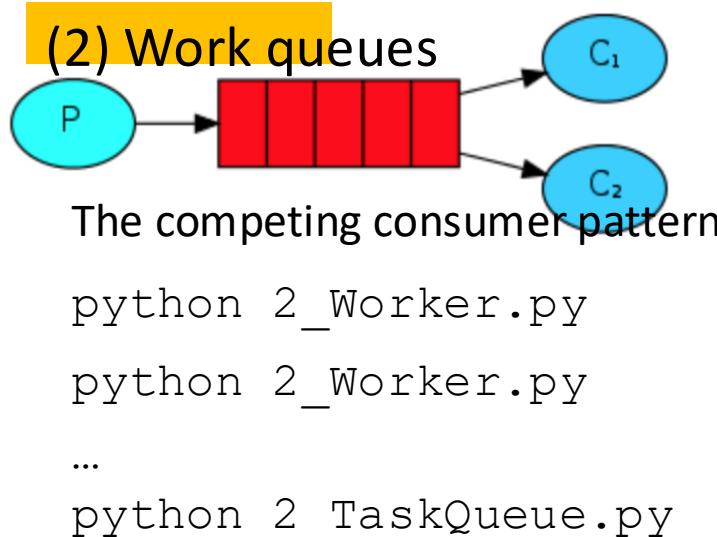
Consumer (C)

Manage acknowledge

```
1  #!/usr/bin/env python
2  import pika
3  import time
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.queue_declare(queue='task_queue', durable=True)
10 print(' [*] Waiting for messages. To exit press CTRL+C')
11
12
13 def callback(ch, method, properties, body):
14     print(" [x] Received %r" % body.decode())
15     time.sleep(body.count(b'.'))
16     print(" [x] Done")
17     ch.basic_ack(delivery_tag=method.delivery_tag)
18
19
20 channel.basic_qos(prefetch_count=1)
21 channel.basic_consume(queue='task_queue', on_message_callback=callback)
22
23 channel.start_consuming()
24
```

To ensure that worker will send ack when completing the task.

(2) Work queue/ Task queue (cont.)



Consumer (C)

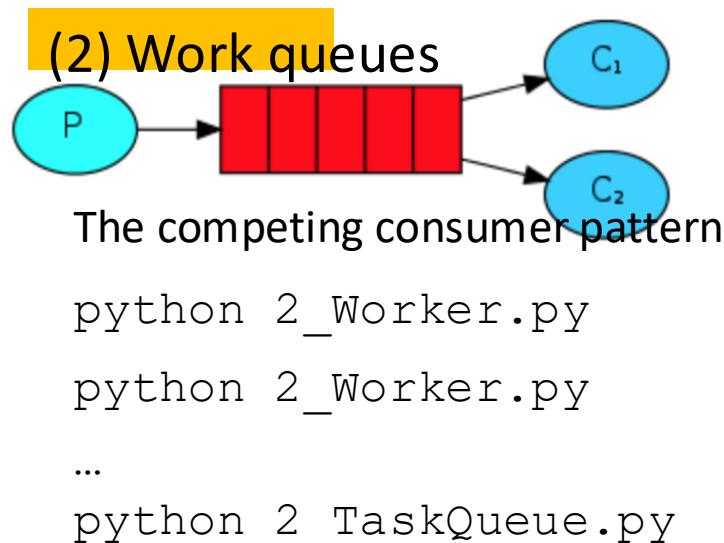
```
1  #!/usr/bin/env python
2  import pika
3  import time
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.queue_declare(queue='task_queue', durable=True)
10 print(' [*] Waiting for messages. To exit press CTRL+C')
11
12
13 def callback(ch, method, properties, body):
14     print(" [x] Received %r" % body.decode())
15     time.sleep(body.count(b'.'))
16     print(" [x] Done")
17     ch.basic_ack(delivery_tag=method.delivery_tag)
18
19
20 channel.basic_qos(prefetch_count=1)
21 channel.basic_consume(queue='task_queue', on_message_callback=callback)
22
23 channel.start_consuming()
24
```

What the difference
between round robin
and fair dispatch?

Fair dispatch

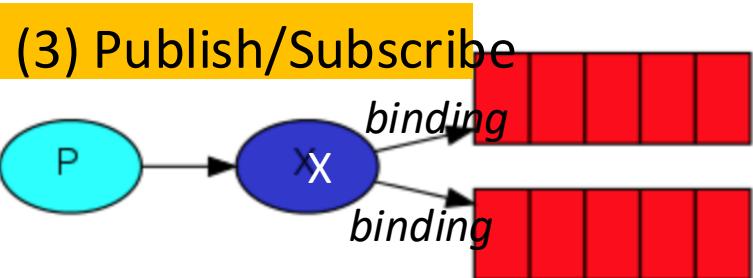
Tell the RabbitMQ not
to dispatch a new task
if the worker is busy.

(2) Work queue/ Task queue (cont.)



DEMO

(3) Publish/ Subscribe (fanout)



Exchange type: direct, topic, headers, **fanout**

```
python 3_Publish_broadcast_log.py  
python 3_Subscribe_log.py
```

RabbitMQ will deliver a message to multiple consumers.

```
docker exec my-rabbitmq-management rabbitmqctl list_exchanges
```

```
docker exec my-rabbitmq-management rabbitmqctl list_bindings
```

Producer (P)

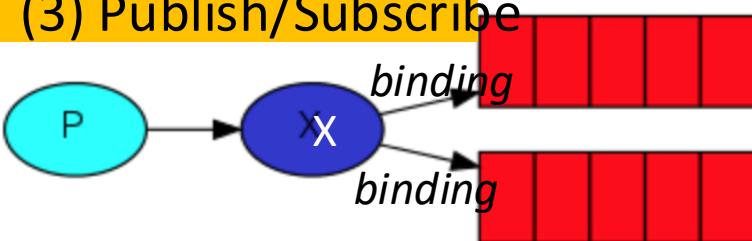
```
1  #!/usr/bin/env python  
2  import pika  
3  import sys  
4  
5  connection = pika.BlockingConnection(  
6      pika.ConnectionParameters(host='localhost'))  
7  channel = connection.channel()  
8  
9  channel.exchange_declare(exchange='logs', exchange_type='fanout')  
10  
11 message = ' '.join(sys.argv[1:]) or "info: Hello World!"  
12 channel.basic_publish(exchange='logs', routing_key='', body=message)  
13 print(" [x] Sent %r" % message)  
14 connection.close()
```

The use of exchange introduces the full messaging model in Rabbit.

Core idea is that a producer never send any messages directly to a queue..

(3) Publish/ Subscribe (fanout) (cont.)

(3) Publish/Subscribe



Exchange type: direct, topic, headers, **fanout**

python 3_Publish_broadcast_log

python 3_Subscribe_log.py

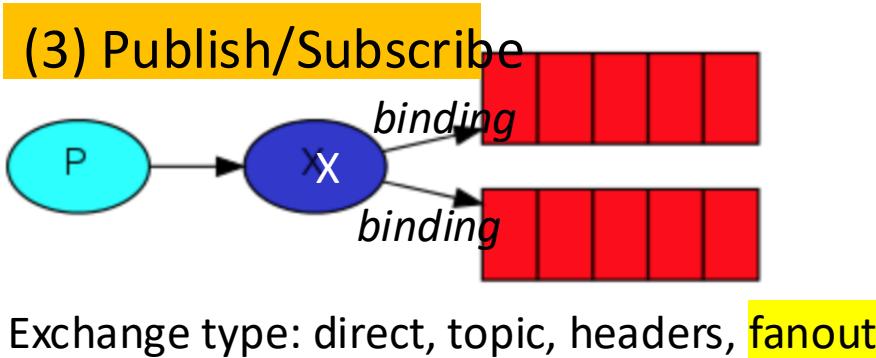
Consumer (C)

```
1  #!/usr/bin/env python
2  import pika
3
4  connection = pika.BlockingConnection(
5      pika.ConnectionParameters(host='localhost'))
6  channel = connection.channel()
7
8  channel.exchange_declare(exchange='logs', exchange_type='fanout')
9
10 result = channel.queue_declare(queue='', exclusive=True)
11 queue_name = result.method.queue
12
13 channel.queue_bind(exchange='logs', queue=queue_name)
14
15 print(' [*] Waiting for logs. To exit press CTRL+C')
16
17 def callback(ch, method, properties, body):
18     print(" [x] %r" % body)
19
20 channel.basic_consume(
21     queue=queue_name, on_message_callback=callback, auto_ack=True)
22
23 Bind exchange with queue
24 channel.start_consuming()
```

We want to hear all messages and only the current flowing messages not the older ones.
ແປລວា?

Let the server choose the random queue name and this queue will be deleted when the consumer's connection is closed.

(3) Publish/ Subscribe (fanout) (cont.)

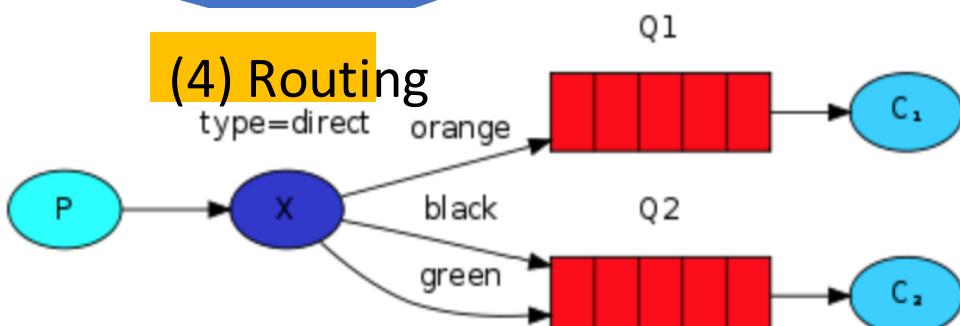


```
python 3_Publish_broadcast_log.py  
python 3_Subscribe_log.py
```

```
python 3_Subscribe_log.py  
python 3_Subscribe_log.py  
python 3_Publish_broadcast_log.py
```

DEMO

(4) Routing



Exchange type: direct, topic, headers, fanout
binding key == routing key

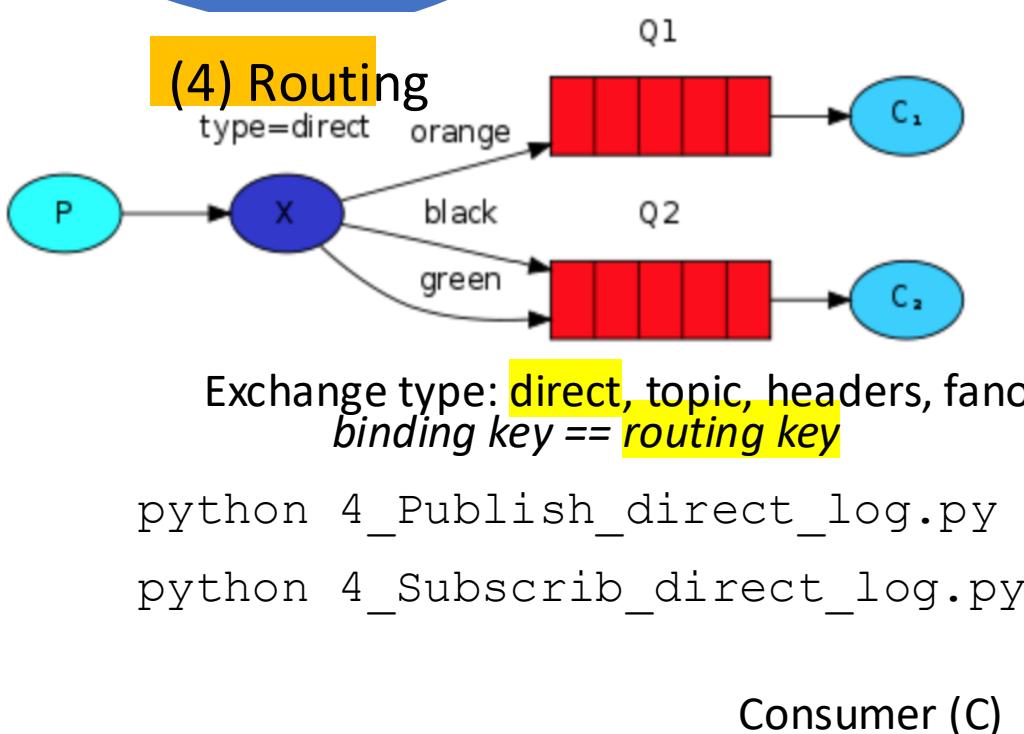
Producer (P)

```
1  #!/usr/bin/env python
2  import pika
3  import sys
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.exchange_declare(exchange='direct_logs', exchange_type='direct')
10
11 severity = sys.argv[1] if len(sys.argv) > 1 else 'info'
12 message = ' '.join(sys.argv[2:]) or 'Hello World!'
13 channel.basic_publish(
14     exchange='direct_logs', routing_key=severity, body=message)
15 print("[x] Sent %r:%r" % (severity, message))
16 connection.close()
```

python 4_Publish_direct_log.py

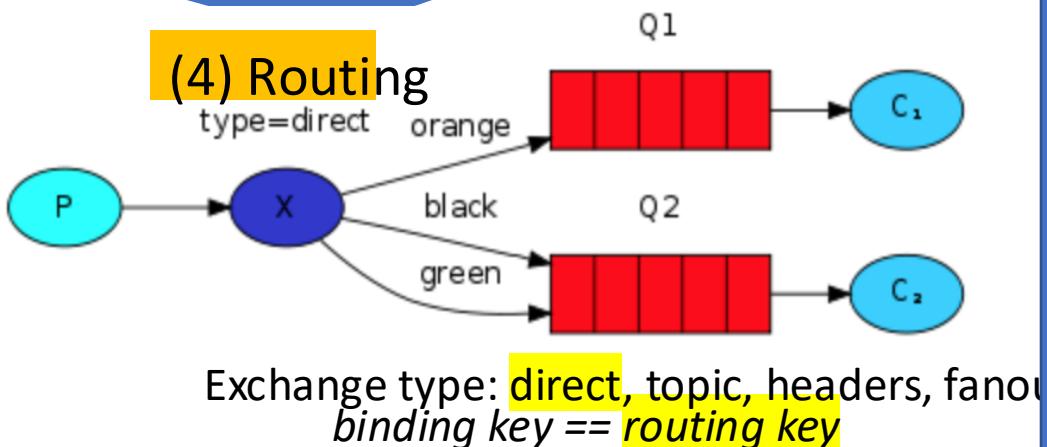
python 4_Subscrib_direct_log.py

(4) Routing



```
1  #!/usr/bin/env python
2  import pika
3  import sys
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.exchange_declare(exchange='direct_logs', exchange_type='direct')
10
11 result = channel.queue_declare(queue='', exclusive=True)
12 queue_name = result.method.queue
13
14 severities = sys.argv[1:]
15 if not severities:
16     sys.stderr.write("Usage: %s [info] [warning] [error]\n" % sys.argv[0])
17     sys.exit(1)
18
19 for severity in severities:
20     channel.queue_bind(
21         exchange='direct_logs', queue=queue_name, routing_key=severity)
22
23 print(' [*] Waiting for logs. To exit press CTRL+C')
24
25
26 def callback(ch, method, properties, body):
27     print(" [x] %r:%r" % (method.routing_key, body))
28
29
30 channel.basic_consume(
31     queue=queue_name, on_message_callback=callback, auto_ack=True)
32
33 channel.start_consuming()
```

(4) Routing



```
python 4_Publish_direct_log.py
```

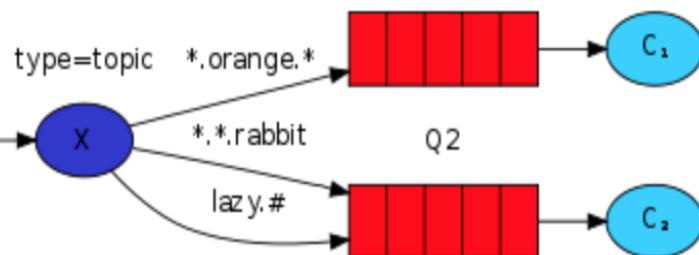
```
python 4_Subscribe_direct_log.py
```

```
python 4_Subscribe_direct_log.py warning error > logs_from_rabbit.log
python 4_Subscribe_direct_log.py info warning error
python 4_Publish_direct_log.py error "GPU card 0 is down!!"
python 4_Publish_direct_log.py info "Add new GPU card!!"
```

DEMO

(5) Topic

(5) Topics
* Can substitute for exactly one word
Can substitute for zero or more words



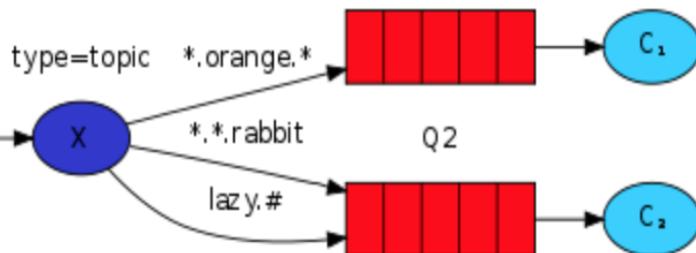
python 5_Publish_topic_log.py

python 5_Subscribe_topic_log.py

```
1  #!/usr/bin/env python
2  import pika
3  import sys
4
5  connection = pika.BlockingConnection(
6      pika.ConnectionParameters(host='localhost'))
7  channel = connection.channel()
8
9  channel.exchange_declare(exchange='topic_logs', exchange_type='topic')
10
11 routing_key = sys.argv[1] if len(sys.argv) > 2 else 'anonymous.info'
12 message = ' '.join(sys.argv[2:]) or 'Hello World!'
13 channel.basic_publish(
14     exchange='topic_logs', routing_key=routing_key, body=message)
15 print("[x] Sent %r:%r" % (routing_key, message))
16 connection.close()
```

(5) Topic

(5) Topics
* Can substitute for exactly one word
Can substitute for zero or more words



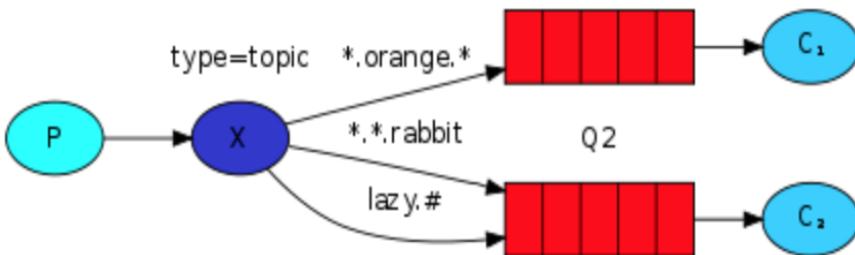
python 5_Publish_topic_log.py

python 5_Subscribe_topic_log.py

```
1 #!/usr/bin/env python
2 import pika
3 import sys
4
5 connection = pika.BlockingConnection(
6     pika.ConnectionParameters(host='localhost'))
7 channel = connection.channel()
8
9 channel.exchange_declare(exchange='topic_logs', exchange_type='topic')
10
11 result = channel.queue_declare('', exclusive=True)
12 queue_name = result.method.queue
13
14 binding_keys = sys.argv[1:]
15 if not binding_keys:
16     sys.stderr.write("Usage: %s [binding_key]...\n" % sys.argv[0])
17     sys.exit(1)
18
19 for binding_key in binding_keys:
20     channel.queue_bind(
21         exchange='topic_logs', queue=queue_name, routing_key=binding_key)
22
23 print(' [*] Waiting for logs. To exit press CTRL+C')
24
25
26 def callback(ch, method, properties, body):
27     print(" [x] %r:%r" % (method.routing_key, body))
28
29
30 channel.basic_consume(
31     queue=queue_name, on_message_callback=callback, auto_ack=True)
32
33 channel.start_consuming()
```

(5) Topic

(5) Topics
* Can substitute for exactly one word
Can substitute for zero or more words



python 5_Publish_topic_log.py

python 5_Subscribe_topic_log.py

```
Terminal 1: python 5_Subscribe_topic_log.py "#"
Terminal 2: python 5_Subscribe_topic_log.py "kern.*"
Terminal 3: python 5_Subscribe_topic_log.py "*.critical"
Terminal 4: python 5_Subscribe_topic_log.py "kern.*" "*critical"
Terminal 5: python 5_Publish_topic_log.py "kern.critical" "kernel error"
```

DEMO

RabbitMQ & Apache Kafka

RabbitMQ	Kafka
Direct messaging: users can set sophisticated rules for message delivery.	Kafka is ideal for handling large amounts of homogeneous messages, such as logs or metric, and it is the right choice for instances with high throughput.
RabbitMQ will be usually used with Cassandra for message playback.	Replay messages
Multiprotocol supports: AMQP, STOMP, MQTT, Web sockets and others.	Big data consideration with e.g., Elastic search, Hadoop
Flexibility: varied point-to-point, request/reply, publish/subscribe	Scaling capability; topics can be split into partitions
Communication: supports both async and sync	Batches: Kafka works best when messages are batched.
Security	Security
Mature platform	Mature platform
Slower than Kafka	Don't come with user-friendly GUI but can be monitored via Kibana

<https://www.upsolver.com/blog/kafka-versus-rabbitmq-architecture-performance-use-case>

<https://dattell.com/data-architecture-blog/kafka-vs-rabbitmq-how-to-choose-an-open-source-message-broker/>

Summary of differences: Kafka vs. RabbitMQ

	RabbitMQ	Kafka
Architecture	RabbitMQ's architecture is designed for complex message routing. It uses the push model. Producers send messages to consumers with different rules.	Kafka uses partition-based design for real-time, high-throughput stream processing. It uses the pull model. Producers publish messages to topics and partitions that consumers subscribe to.
Message handling	RabbitMQ brokers monitor message consumption. It deletes messages after they're consumed. It supports message priorities.	Consumers keep track of message retrieval with an offset tracker. Kafka retains messages according to the retention policy. There's no message priority.
Performance	RabbitMQ has low latency. It sends thousands of messages per second.	Kafka has real-time transmission of up to millions of messages per second.
Programming language and protocol	RabbitMQ supports a broad range of languages and legacy protocols.	Kafka has limited choices of programming languages. It uses binary protocol over TCP for data transmission.

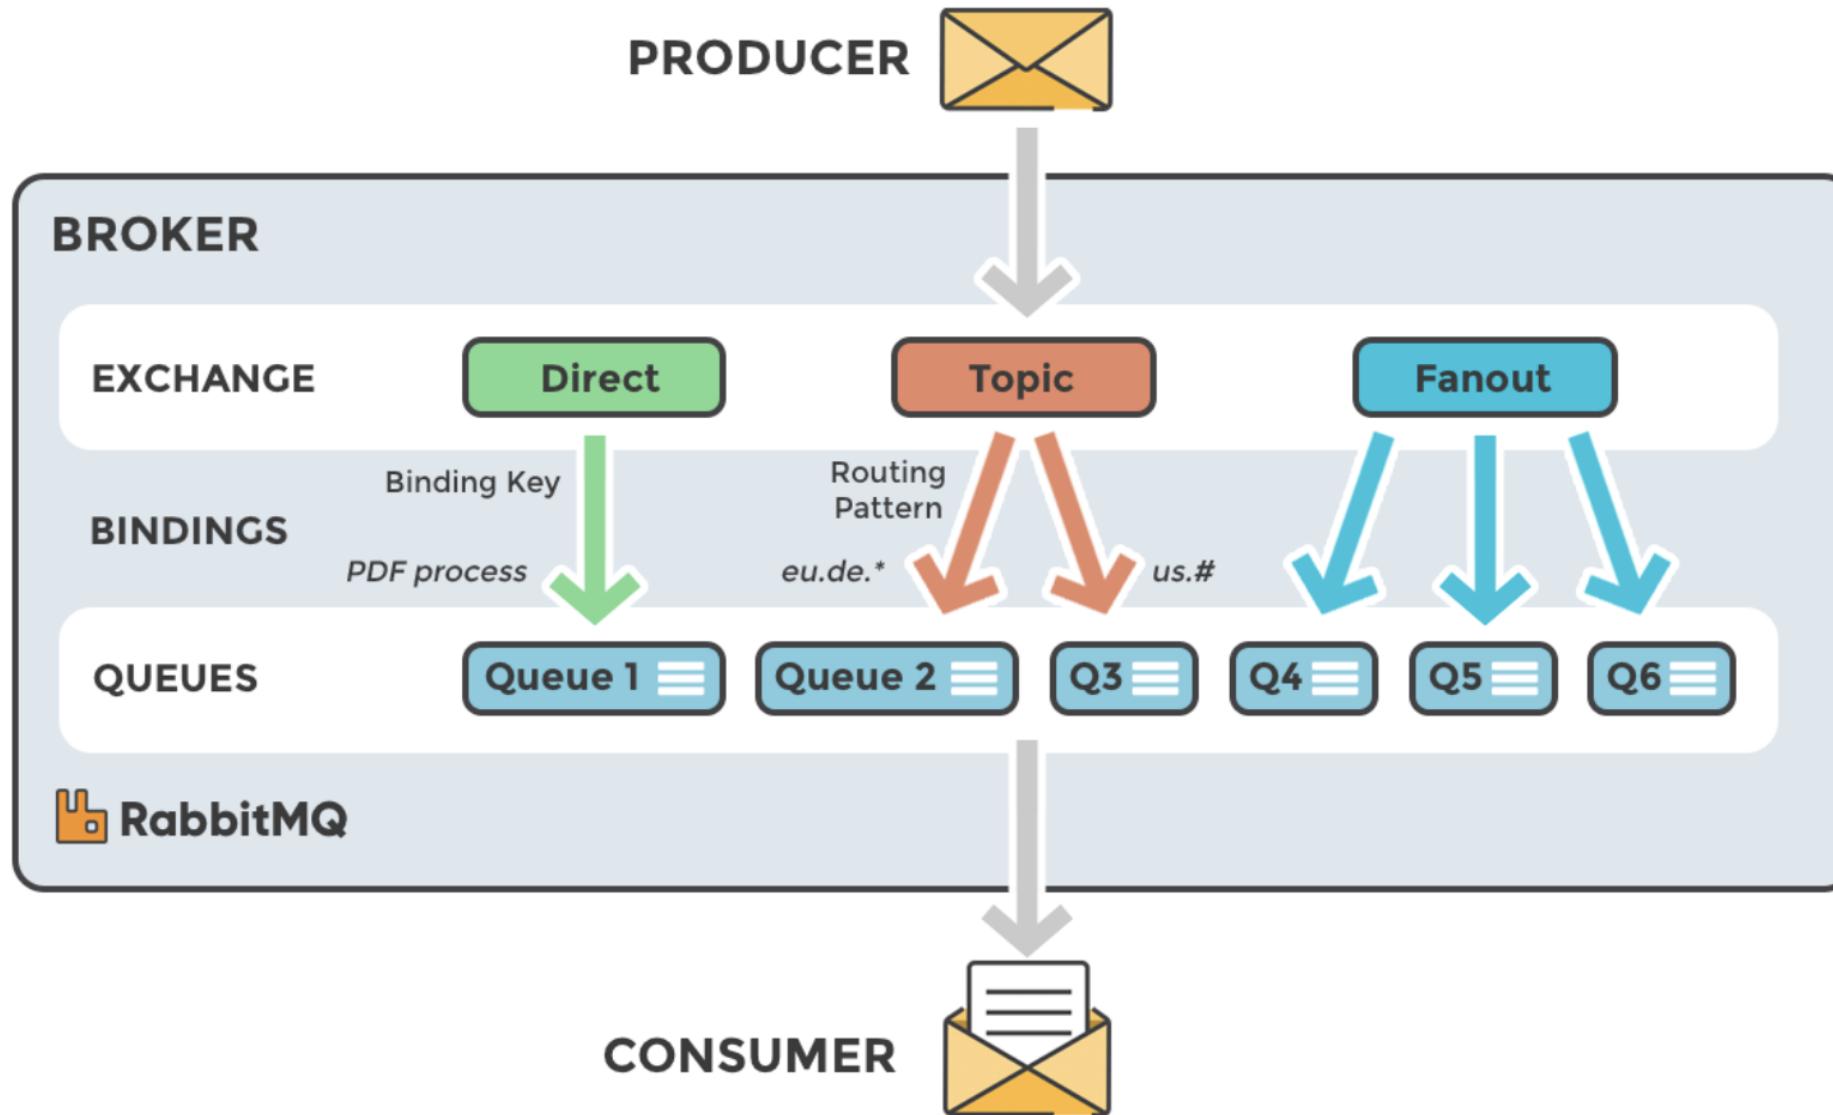
RabbitMQ & Apache Kafka



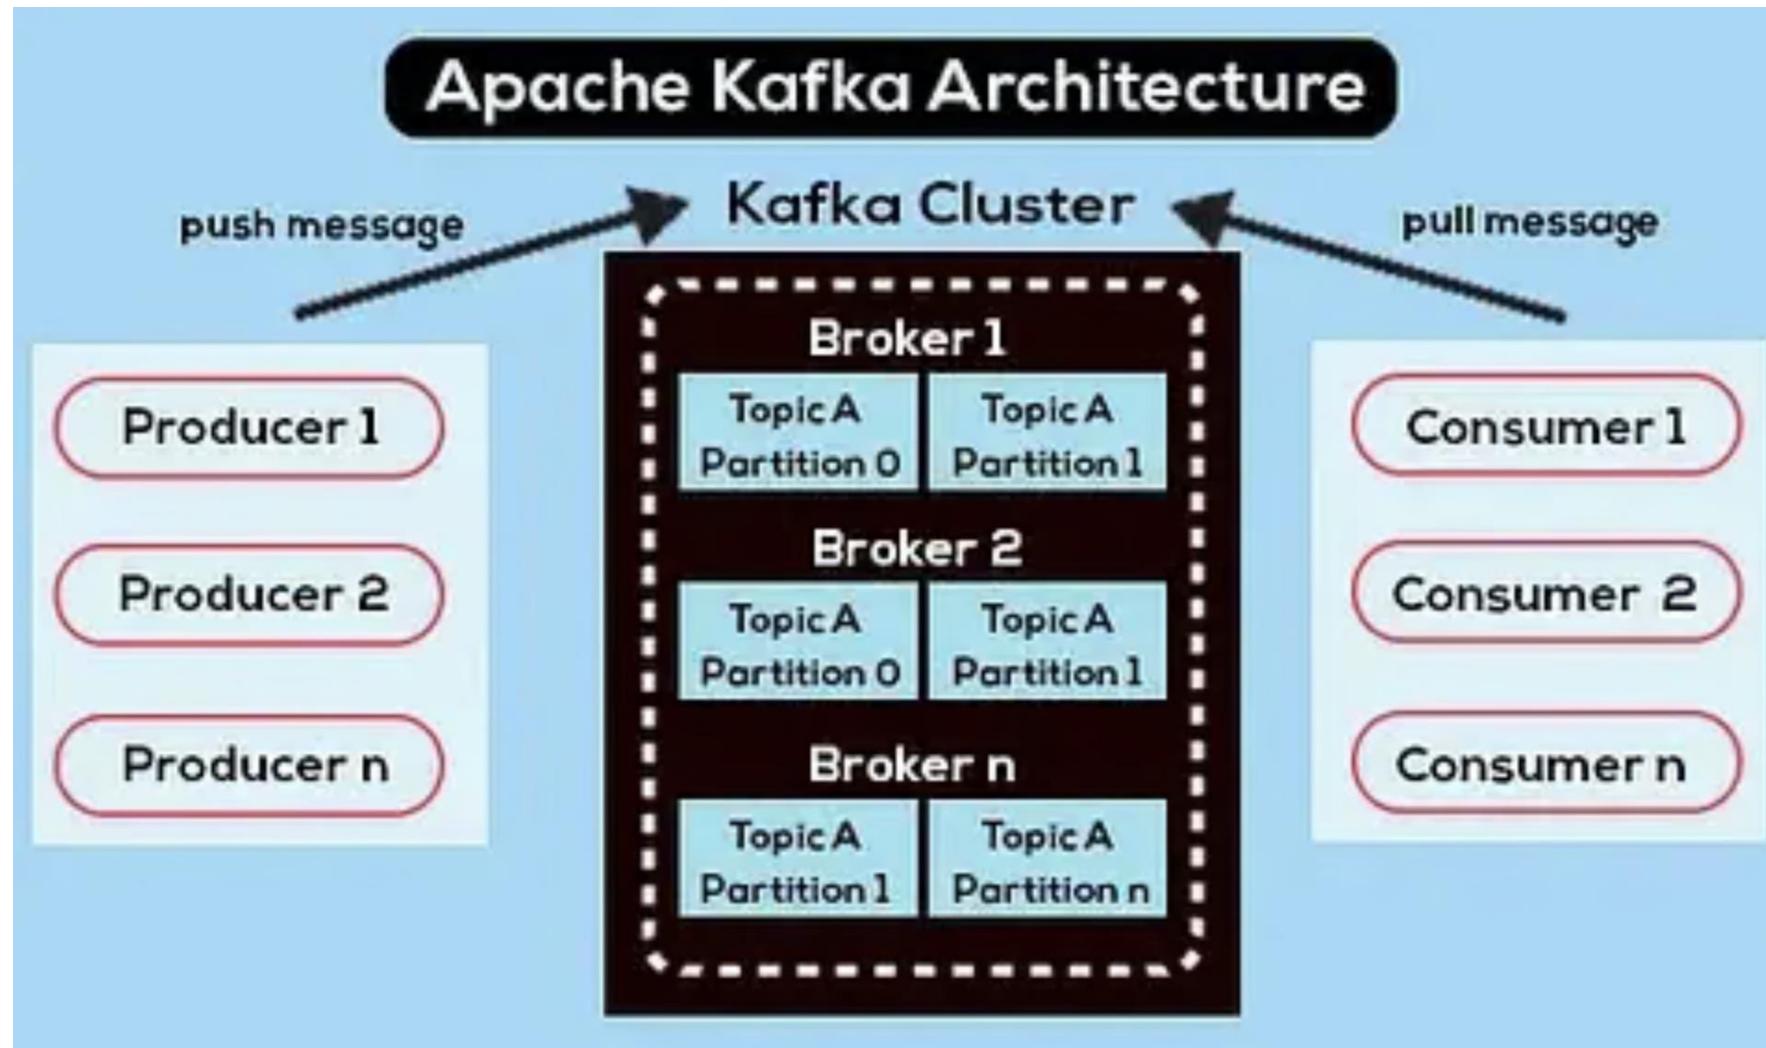
RabbitMQ: Used for high throughput and reliable background jobs, integration and intercommunication between and within applications, perform complex routing to consumers, integrate multiple applications and services with non-trivial routing logic.

Kafka: Best used for basic streaming without complex routing with maximum throughput, ideal for event-sourcing, stream processing, multi-stage pipelines, routinely audited systems, real-time processing and analyzing data

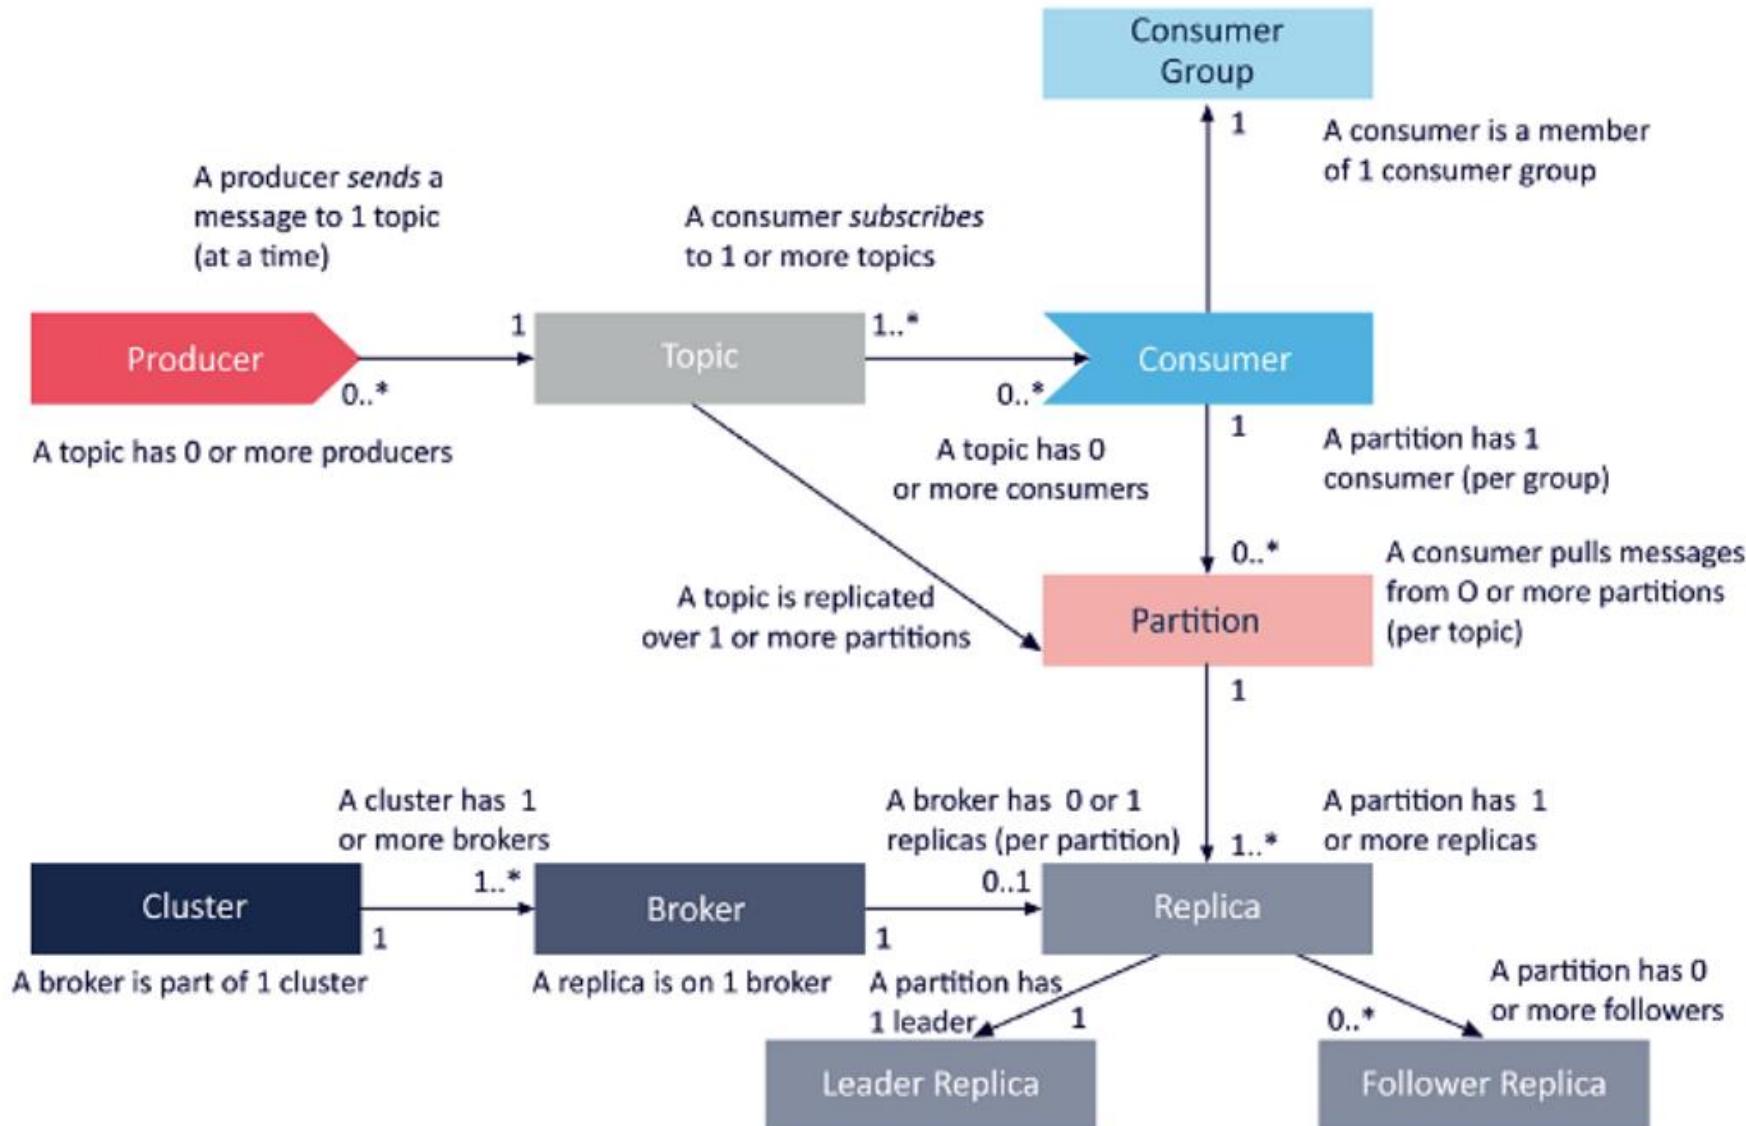
RabbitMQ Overall Architecture



Kafka Overall Architecture



Main Components in Kafka



Assignment

- Go to get starter source codes from

<https://drive.google.com/drive/folders/1iRNlwWf8uSkfp3TriQm9QZCHjgx75CEk?usp=sharing>

- Try running Kafka as shown in the starter kit (how_to_start_kafka.txt)
- Based on the given RabbitMQ and Kafka codes, compare the performance of RabbitMQ and Kafka based on round-robin work, where a producer produces 100K messages with the four consumers and the varied message sizes of 0.1, 0.5, and 1KB.
- Things to submit
 - 1) Your modified codes with the experiments
 - 2) Three plots (a message size per plot) comparing the throughputs between the two brokers.

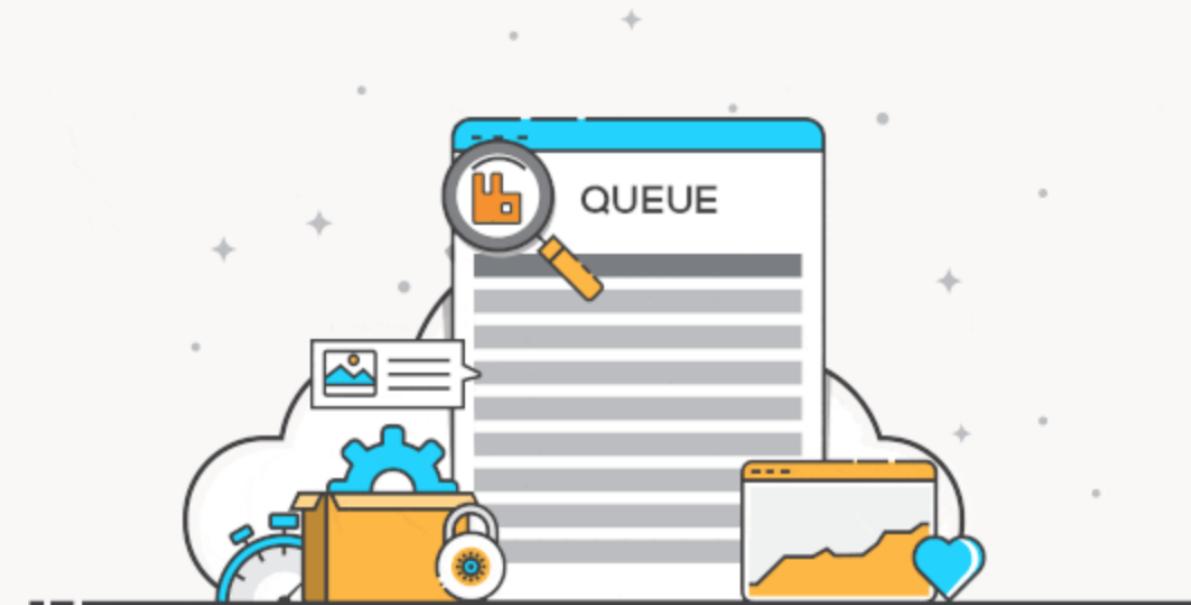
RabbitMQ as a Service

+ .
o



RabbitMQ as a Service

Managing the largest fleet of RabbitMQ clusters in the world



Get a managed RabbitMQ server today